

Compiling and Linking and Interfacing Multiple Programming Languages

Dr. Axel Kohlmeyer

Research Professor, Dept. of Mathematics
Associate Director, ICMS
College of Science and Technology
Temple University, Philadelphia

<http://sites.google.com/site/akohlmey/>

a.kohlmeyer@temple.edu

Pre-process / Compile / Link

- Creating an executable includes multiple steps
- The “compiler” (gcc) is a wrapper for several commands that are executed in succession
- The “compiler flags” similarly fall into categories and are handed down to the respective tools
- The “wrapper” selects the compiler language from source file name, but links “its” runtime
- We will look into a C example first, since this is the language the OS is (mostly) written in

A simple C Example

- Consider the minimal C program 'hello.c':
#include <stdio.h>
int main(int argc, char **argv)
{
 printf("hello world\n");
 return 0;
}
- i.e.: what happens, if we do:
 > gcc -o hello hello.c
 (tr try: **gcc -v -o hello hello.c**)

Step 1: Pre-processing

- Pre-processing is mandatory in C (and C++)
- Pre-processing will handle '#' directives
 - File inclusion with support for nested inclusion
 - Conditional compilation and Macro expansion
- In this case: **/usr/include/stdio.h**
 - and all files are included by it - are inserted and the contained macros expanded
- Use -E flag to stop after pre-processing:
> **cc -E -o hello.pp.c hello.c**

Step 2: Compilation

- Compiler converts a high-level language into the specific instruction set of the target CPU
- Individual steps:
 - Parse text (lexical + syntactical analysis)
 - Do language specific transformations
 - Translate to internal representation units (IRs)
 - Optimization (reorder, merge, eliminate)
 - Replace IRs with pieces of assembler language
- Try:> **gcc -S hello.c** (produces **hello.s**)

Compilation cont'd

```
.file "hello.c"
.section .rodata
.LC0:
.string "hello, world!"
.text
.globl main
.type main, @function
main:
    pushl    %ebp
    movl     %esp, %ebp
    andl     $-16, %esp
    subl     $16, %esp
    movl     $.LC0, (%esp)
    call     puts
    movl     $0, %eax
    leave
    ret
.size      main, .-main
.ident     "GCC: (GNU) 4.5.1 20100924 (Red Hat 4.5.1-4)"
.section   .note.GNU-stack,"",@progbits
```

gcc replaced printf with puts

try: gcc -fno-builtin -S hello.c

```
#include <stdio.h>
int main(int argc,
          char **argv)
{
    printf("hello world\n");
    return 0;
}
```


Step 3: Assembler / Step 4: Linker

- Assembler (as) translates assembly to binary
 - Creates so-called object files (in ELF format)

```
Try: > gcc -c hello.c
```

```
Try: > nm hello.o
```

```
000000000 T main
```

```
          U puts
```

- Linker (ld) puts binary together with startup code and required libraries
- Final step, result is executable.

```
Try: > gcc -o hello hello.o
```

Adding Libraries

- Example 2: exp.c

```
#include <math.h>
#include <stdio.h>
int main(int argc, char **argv)
{
    double a=2.0;
    printf("exp(2.0)=%f\n", exp(a));
    return 0;
}
```

- > gcc -o exp exp.c
Fails with “undefined reference to 'exp'”. Add: -lm
- > gcc -O3 -o exp exp.c
Works due to inlining at high optimization level.

Symbols in Object Files & Visibility

- Compiled object files have multiple sections and a symbol table describing their entries:
 - “Text”: this is executable code
 - “Data”: pre-allocated variables storage
 - “Constants”: read-only data
 - “Undefined”: symbols that are used but not defined
 - “Debug”: debugger information (e.g. line numbers)
- Entries in the object files can be inspected with either the “nm” tool or the “readelf” command

Example File: visibility.c

```
static const int val1 = -5;
const int val2 = 10;
static int val3 = -20;
int val4 = -15;
extern int errno;
```

```
static int add_abs(const int v1, const int v2) {
    return abs(v1)+abs(v2);
}
```

```
int main(int argc, char **argv) {
    int val5 = 20;
    printf("%d / %d / %d\n",
        add_abs(val1,val2),
        add_abs(val3,val4),
        add_abs(val1,val5));
    return 0;
}
```

```
nm visibility.o:
00000000 t add_abs
                U errno
00000024 T main
                U printf
00000000 r val1
00000004 R val2
00000000 d val3
00000004 D val4
```

What Happens During Linking?

- Historically, the linker combines a “startup object” (crt1.o) with all compiled or listed object files, the C library (libc) and a “finish object” (crtn.o) into an executable (a.out)
- With current compilers it is more complicated
- The linker then “builds” the executable by matching undefined references with available entries in the symbol tables of the objects
- crt1.o has an undefined reference to “main” thus C programs start at the main() function

Static Libraries

- Static libraries built with the “ar” command are collections of objects with a global symbol table
- When linking to a static library, object code is copied into the resulting executable and all direct addresses recomputed (e.g. for “jumps”)
- Symbols are resolved “from left to right”, so circular dependencies require to list libraries multiple times or use a special linker flag
- When linking only the name of the symbol is checked, not whether its argument list matches

Shared Libraries

- Shared libraries are more like executables that are missing the `main()` function
- When linking to a shared library, a marker is added to load the library by its “generic” name (soname) and the list of undefined symbols
- When resolving a symbol (function) from shared library all addresses have to be recomputed (relocated) on the fly.
- The shared linker program is executed first and then loads the executable and its dependencies

Differences When Linking

- Static libraries are fully resolved “left to right”; circular dependencies are only resolved between explicit objects or inside a library
-> need to specify libraries multiple times
or use: **-Wl,--start-group (...) -Wl,--end-group**
- Shared libraries symbols are not fully resolved at link time, only checked for symbols required by the object files. Full check only at runtime.
- Shared libraries may depend on other shared libraries whose symbols will be globally visible

Dynamic Linker Properties

- Linux defaults to dynamic libraries:
> ldd hello
linux-gate.so.1 => (0x0049d000)
libc.so.6 => /lib/libc.so.6
(0x005a0000)
/lib/ld-linux.so.2 (0x0057b000)
- **/etc/ld.so.conf, LD_LIBRARY_PATH** define where to search for shared libraries
- **gcc -Wl, -rpath, /some/dir** will encode **/some/dir** into the binary for searching

Using LD_PRELOAD

- Using the LD_PRELOAD environment variable, symbols from a shared object can be preloaded into the global object table and will override those in later resolved shared libraries
=> replace specific functions in a shared library
- Example: override log() with a faster version:

```
#include "amdlibm.h"  
double log(double x) { return amd_log(x); }  
gcc -shared -o fasterlog.so faster.c -lamdlibm
```
- LD_PRELOAD=./fasterlog.so ./myprog-with

Before LD_PRELOAD

PerfTop: 8016 irqs/sec kernel: 9.9% exact: 0.0% [1000Hz cycles], (all, 8 CPUs)

| samples | pcnt | function | DSO |
|----------|-------|------------------------|--------------------------------------|
| 53462.00 | 52.2% | __ieee754_log | /lib64/libm-2.12.so |
| 10490.00 | 10.3% | R_binary | /opt/bin/R-2.13.0/lib64/R/bin/exec/R |
| 8704.00 | 8.5% | clear_page_c | [kernel.kallsyms] |
| 5737.00 | 5.6% | __ieee754_exp | /lib64/libm-2.12.so |
| 4645.00 | 4.5% | math1 | /opt/bin/R-2.13.0/lib64/R/bin/exec/R |
| 3070.00 | 3.0% | __log | /lib64/libm-2.12.so |
| 3020.00 | 3.0% | __isnan | /lib64/libc-2.12.so |
| 2094.00 | 2.0% | R_gc_internal | /opt/bin/R-2.13.0/lib64/R/bin/exec/R |
| 1643.00 | 1.6% | do_summary | /opt/bin/R-2.13.0/lib64/R/bin/exec/R |
| 1251.00 | 1.2% | __isnan@plt | /opt/bin/R-2.13.0/lib64/R/bin/exec/R |
| 1210.00 | 1.2% | real_relop | /opt/bin/R-2.13.0/lib64/R/bin/exec/R |
| 1161.00 | 1.1% | __GI__exp | /lib64/libm-2.12.so |
| 754.00 | 0.7% | __isnan | /lib64/libm-2.12.so |
| 739.00 | 0.7% | R_log | /opt/bin/R-2.13.0/lib64/R/bin/exec/R |
| 553.00 | 0.5% | __kernel_standard | /lib64/libm-2.12.so |
| 550.00 | 0.5% | do_abs | /opt/bin/R-2.13.0/lib64/R/bin/exec/R |
| 462.00 | 0.5% | __mul | /lib64/libm-2.12.so |
| 439.00 | 0.4% | coerceToReal | /opt/bin/R-2.13.0/lib64/R/bin/exec/R |
| 413.00 | 0.4% | finite | /lib64/libm-2.12.so |
| 358.00 | 0.3% | log@plt | /opt/bin/R-2.13.0/lib64/R/bin/exec/R |
| 182.00 | 0.2% | get_page_from_freelist | [kernel.kallsyms] |
| 120.00 | 0.1% | __alloc_pages_nodemask | [kernel.kallsyms] |

After LD_PRELOAD

PerfTop: 8020 irqs/sec kernel:17.2% exact: 0.0% [1000Hz cycles], (all, 8 CPUs)

| samples | pcnt | function | DSO |
|----------|-------|-------------------|--|
| 24702.00 | 19.5% | __amd_bas64_log | /opt/libs/fastermath-0.1/libamdlibm.so |
| 22270.00 | 17.6% | R_binary | /opt/binf/R-2.13.0/lib64/R/bin/exec/R |
| 18463.00 | 14.6% | clear_page_c | [kernel.kallsyms] |
| 10480.00 | 8.3% | __ieee754_exp | /lib64/libm-2.12.so |
| 9834.00 | 7.8% | math1 | /opt/binf/R-2.13.0/lib64/R/bin/exec/R |
| 9155.00 | 7.2% | log | /opt/libs/fastermath-0.1/fasterlog.so |
| 6269.00 | 5.0% | __isnan | /lib64/libc-2.12.so |
| 4214.00 | 3.3% | R_gc_internal | /opt/binf/R-2.13.0/lib64/R/bin/exec/R |
| 3074.00 | 2.4% | do_summary | /opt/binf/R-2.13.0/lib64/R/bin/exec/R |
| 2285.00 | 1.8% | real_relop | /opt/binf/R-2.13.0/lib64/R/bin/exec/R |
| 2257.00 | 1.8% | __isnan@plt | /opt/binf/R-2.13.0/lib64/R/bin/exec/R |
| 2076.00 | 1.6% | __GI__exp | /lib64/libm-2.12.so |
| 1346.00 | 1.1% | R_log | /opt/binf/R-2.13.0/lib64/R/bin/exec/R |
| 1213.00 | 1.0% | do_abs | /opt/binf/R-2.13.0/lib64/R/bin/exec/R |
| 1075.00 | 0.8% | __kernel_standard | /lib64/libm-2.12.so |
| 894.00 | 0.7% | coerceToReal | /opt/binf/R-2.13.0/lib64/R/bin/exec/R |
| 780.00 | 0.6% | __mul | /lib64/libm-2.12.so |
| 756.00 | 0.6% | finite | /lib64/libm-2.12.so |
| 729.00 | 0.6% | amd_log@plt | /opt/libs/fastermath-0.1/fasterlog.so |
| 706.00 | 0.6% | amd_log | /opt/libs/fastermath-0.1/libamdlibm.so |
| 674.00 | 0.5% | log@plt | /opt/binf/R-2.13.0/lib64/R/bin/exec/R |

Difference Between C and Fortran

- Basic compilation principles are the same
=> preprocess, compile, assemble, link
- In Fortran, symbols are case insensitive
=> most compilers translate them to lower case
- In Fortran symbol names may be modified to make them different from C symbols
(e.g. append one or more underscores)
- Fortran entry point is not “main” (no arguments)
PROGRAM => MAIN__ (in gfortran)
- C-like main() provided as startup (to store args)

Pre-processing in C and Fortran

- Pre-processing is mandatory in C/C++
- Pre-processing is optional in Fortran
- Fortran pre-processing enabled implicitly via file name: name.F, name.F90, name.FOR
- Legacy Fortran packages often use /lib/cpp:
/lib/cpp -C -P -traditional -o name.f name.F
 - -C : keep comments (may be legal Fortran code)
 - -P : no '#line' markers (not legal Fortran syntax)
 - -traditional : don't collapse whitespace (incompatible with fixed format sources)

Fortran Symbols Example

| | |
|-------------------------|--------------------------------|
| SUBROUTINE GREET | 0000006d t MAIN__ |
| PRINT*, 'HELLO, WORLD!' | U _gfortran_set_args |
| END SUBROUTINE GREET | U _gfortran_set_options |
| | U _gfortran_st_write |
| program hello | U _gfortran_st_write_done |
| call greet | U _gfortran_transfer_character |
| end program | 00000000 T greet__ |
| | 0000007a T main |

- “program” becomes symbol “MAIN__” (compiler dependent)
- “subroutine” name becomes lower case with '_' appended
- several “undefines” with '_gfortran' prefix
 - => calls into the Fortran runtime library, libgfortran
- cannot link object with “gcc” alone, need to add -lgfortran
 - => cannot mix and match Fortran objects from different compilers

Fortran 90+ Modules

- When subroutines or variables are defined inside a module, they have to be hidden

```
module func
  integer :: val5, val6
contains
  integer function add_abs(v1,v2)
    integer, intent(in) :: v1, v2
    add_abs = iabs(v1)+iabs(v2)
  end function add_abs
end module func
```

- gfortran creates the following symbols:

```
000000000 T __func_MOD_add_abs
000000000 B __func_MOD_val5
000000004 B __func_MOD_val6
```

The Next Level: C++

- In C++ functions with different number or type of arguments can be defined (overloading)
=> encode prototype into symbol name:

Example : symbol for `int add_abs(int,int)` becomes: `_ZL7add_absii`

- Note: the return type is not encoded
- C++ symbols are no longer compatible with C
=> add 'extern "C"' qualifier for C style symbols
- C++ symbol encoding is compiler specific

C++ Namespaces and Classes vs. Fortran 90 Modules

- Fortran 90 modules share functionality with classes and namespaces in C++
- C++ namespaces are encoded in symbols
Example: `int func::add_abs(int,int)`
becomes: `_ZN4funcL7add_absEii`
- C++ classes are encoded the same way
- Figuring out which symbol to encode into the object as undefined is the job of the compiler
- When using the gdb debugger use '::' syntax

Why We Need Header or Module Files

- The linker is “blind” for any language specific properties of a symbol => checking of the validity of the interface of a function is only possible during compilation
- A header or module file contains the prototype of the function (not the implementation) and the compiler can compare it to its use
- Important: header/module has to match library
=> Problem with FFTW-2.x: cannot tell if library was compiled for single or double precision

Calling C from Fortran 77

- Need to make C function look like Fortran 77
 - Append underscore (except on AIX, HP-UX)
 - Call by reference conventions
 - Best only used for “subroutine” constructs (cf. MPI) as passing return value of functions varies a lot:

```
void add_abs_(int *v1,int *v2,int *res){  
    *res = abs(*v1)+abs(*v2);}
```
- Arrays are always passed as “flat” 1d arrays by providing a pointer to the first array element
- Strings are tricky (no terminal 0, length added)

Calling C from Fortran 77 Example

```
void sum_abs_(int *in, int *num, int *out) {  
    int i, sum;  
    sum = 0;  
    for (i=0; i < *num; ++i) { sum += abs(in[i]);}  
    *out = sum;  
    return;  
}
```

```
/* fortran code:  
integer, parameter :: n=200  
integer :: s, data(n)  
  
call SUM_ABS(data, n, s)  
print*, s  
*/
```

Calling Fortran 77 from C

- Inverse from previous, i.e. need to add underscore and use lower case (usually)
- Difficult for anything but Fortran 77 style calls since Fortran 90+ features need extra info
 - Shaped arrays, optional parameters, modules
- Arrays need to be “flat”,
C-style multi-dimensional arrays are lists of pointers to individual pieces of storage, which may not be consecutive
=> use 1d and compute position

Calling Fortran 77 From C Example

```
subroutine sum_abs(in, num, out)
  integer, intent(in)  :: num, in(num)
  integer, intent(out) :: out
  Integer               :: i, sum
  sum = 0
  do i=1,num
    sum = sum + ABS(in(i))
  end do
  out = sum
end subroutine sum_abs
!! c code:
!   const int n=200;
!   int data[n], s;
!   sum_abs_(data, &n, &s);
!   printf("%d\n", s);
```

Modern Fortran vs C Interoperability

- Fortran 2003 introduces a standardized way to tell Fortran how C functions look like and how to make Fortran functions have a C-style ABI
- Module “iso_c_binding” provides kind definition: e.g. C_INT, C_FLOAT, C_SIGNED_CHAR
- Subroutines can be declared with “BIND(C)”
- Arguments can be given the property “VALUE” to indicate C-style call-by-value conventions
- String passing tricky, needs explicit 0-terminus

Calling C from Fortran 03 Example

```
int sum_abs(int *in, int num) {
    int i,sum;
    for (i=0,sum=0;i<num;++i) {sum += abs(in[i]);}
    return sum;
}
/* fortran code:
use iso_c_binding, only: c_int
interface
    integer(c_int) function sum_abs(in, num) bind(C)
        use iso_c_binding, only: c_int
        integer(c_int), intent(in) :: in(*)
        integer(c_int), value :: num
    end function sum_abs
end interface
integer(c_int), parameter :: n=200
integer(c_int) :: data(n)
print*, SUM_ABS(data,n) */
```

Calling Fortran 03 From C Example

```
subroutine sum_abs(in, num, out) bind(c)
  use iso_c_binding, only : c_int
  integer(c_int), intent(in)  :: num, in(num)
  integer(c_int), intent(out) :: out
  integer(c_int),              :: i, sum
  sum = 0
  do i=1,num
    sum = sum + ABS(in(i))
  end do
  out = sum
end subroutine sum_abs
```

```
!! c code:
!   const int n=200;
!   int data[n], s;
!   sum_abs(data, &n, &s);
!   printf("%d\n", s);
```


Linking Multi-Language Binaries

- Inter-language calls via mutual C interface only due to name “mangling” of C++ / Fortran 90+
=> extern “C”, ISO_C_BINDING, C wrappers
- Fortran “main” requires Fortran compiler for link
- Global static C++ objects require C++ for link
=> avoid static objects (good idea in general)
- Either language requires its runtime for link
=> GNU: -lstdc++ and -lgfortran
=> Intel: “its complicated” (use -# to find out)
more may be needed (-lgomp, -lpthread, -lm)

Dynamic Linking via dlopen()

- POSIX compliant C libraries allow loading of shared objects at runtime via dlopen()/dlsym()
 - Calls to dlopen() open a handle to shared object; lookup of this file is subject to same rules as dynamic library searches
 - Calls to dlsym() look up symbol by its name in shared object pointed to by handle; returns pointer; for functions need to cast/assign to function pointer
 - Calls to dlclose() unload shared object (if last user) and revoke assignments to code made by dlsym()

Example: static program test-0.c

```
#include <stdio.h>

void hello()
    puts("Hello, World");
}

int main(int argc, char **argv)
{
    void (*hi)(); /* function pointer variable */

    hi = &hello; /* initialize function pointer */

    (*hi)();      /* this is the same as: hello(); */
    return 0;
}

/* compile with: gcc -o test-0 -Wall -O test-0.c */
```

Example: main program test-1.c

```
#include <dlfcn.h>

int main(int argc, char **argv)
{
    void *handle;          /* handle for dynamic object */
    void (*hi)();          /* function pointer for symbol */

    handle = dlopen("./hello.so", RTLD_LAZY);
    if (handle) {
        hi = (void (*)()) dlsym(handle, "hello");
        (*hi)();
        dlclose(handle);
    }
    return 0;
}

/* compile with: gcc -o test-1 -Wall -O test-1.c -ldl
   add -rdynamic if shared object needs symbols in main */
```

Example: shared object hello.c

```
#include<stdio.h>
```

```
void hello(void)
```

```
{
```

```
    puts("Hello, World!");
```

```
}
```

```
/*
```

```
compile: gcc -shared -o hello.so -fPIC -Wall -O hello.c
```

```
*/
```

- With this setup, `hello.c` can be changed and `hello.so` recompiled without having to recompile and re-link `test-1`.
- Thus access to `test-1.c` is not needed.

Extending Python with ctypes

- The ctypes module in python provides an interface to dlopen()/dlsym() and thus allows to call compiled C code from python.
- Support for dll files on Windows is also included
- Since symbols in compiled objects have no information about calling sequence and return values, this has to be set on the python side
- Incorrect use can lead to segmentation faults or corrupted data; often prototypes are needed

Example: calling hello.c from python

```
#!/usr/bin/env python

from ctypes import *

# import shared object on POSIX compatible OS
dso = CDLL("./hello.so")

# call symbol in shared object as function w/o args
dso.hello()
```

- This python script does pretty much the same thing as the `test-1` compiled program
- Since there are no arguments and no return values, no code needs to know about the other

Arguments & Return Value

- By default ctypes will assume arguments and return values are standard size integer

```
#include<stdio.h>
```

```
int sum_of_int(int a, int b) {  
    int c = a + b;  
    printf("sum of %d and %d is %d\n",a,b,c);  
    return c;  
}
```

```
-----  
#!/usr/bin/env python  
from ctypes import *  
dso = CDLL("./sum.so")  
isum = dso.sum_of_int(1,2)  
print "Integer sum is: ", isum
```

Prototypes with ctypes

- If argument and/or return value are of different type, ctypes needs to be informed about it; works similar to prototypes in C

```
#!/usr/bin/env python
from ctypes import *
dso = CDLL("./sum.so")
dso.sum_of_int.argtypes = [ c_int, c_int ]
dso.sum_of_int.restype = c_int
isum = dso.sum_of_int(1,2)
print ("Integer sum w/ prototypes is: ", isum)

dso.sum_of_double.argtypes = [ c_double, c_double ]
dso.sum_of_double.restype = c_double
dsum = dso.sum_of_double(0.5,2.5)
print ("Double sum w/ prototypes is: ", dsum)
```

Passing Strings

- Strings in python are read-only, thus when a C-function will modify a string we have to use `create_string_buffer()`

```
#!/usr/bin/env python
from ctypes import *
dso = CDLL("./hello.so")
```

```
# hello() in hello.so takes a "char *" argument
dso.hello.argtypes = [ c_char_p ]
dso.hello(b"World")
```

```
# create buffer for mutable string data
buf = create_string_buffer(b"World")
dso.hello(buf)
```


Passing Arrays

- When passing allocatable objects like arrays, it is usually best to do the allocating in python. ctypes offers constructors for all basic types

```
#!/usr/bin/env python
from ctypes import *
dso = CDLL("./sum.so")
num = 10
dlist = (c_double * num)() # (primitive * length)()
for i in range(num):
    dlist[i] = 0.333*(i*0.5)
# note the use of POINTER()
dso.sum_of_doubles.argtypes=[POINTER(c_double),c_int]
dso.sum_of_doubles.restype = c_double
dsum = dso.sum_of_doubles(dlist,num)
print ("Double sum is: ", dsum)
```

Passing Structs

- Even complex storage elements like struct can be managed by ctypes. Derive a class from Structure that mimics the corresponding C-type

```
#!/usr/bin/env python
from ctypes import *
dso = CDLL("./data.so")

class parm(Structure):
    _fields_ = [ ("type", c_int), ("label", c_char_p),
                 ("epsilon", c_double), ("sigma", c_double) ]
# use constructor to initialize struct
p = parm(type=1, label=b"LJ-12-6", epsilon=0.1, sigma=3.4)
# p is passed by value, to pass by reference use byref(p)
dso.pass_by_value(p)
dso.pass_by_reference(byref(p))
```

Passing Structs (2)

- Below is the corresponding C code:

```
#include<stdio.h>

struct parm { int  type; char *label;
              double epsilon, sigma;
};

void pass_by_value(struct parm p) {
    printf("type=%d  label=%s epsilon=%g  sigma=%g\n",
          p.type, p.label, p.epsilon, p.sigma);
}

void pass_by_reference(struct parm *p) {
    printf("type=%d  label=%s epsilon=%g  sigma=%g\n",
          p->type, p->label, p->epsilon, p->sigma);
}
```

Interfacing Fortran with f2py

- Interfacing Fortran with python is both easier and more complicated than interfacing C
 - The Fortran ABI can be much more complex and is more compiler specific than the C ABI
 - The numpy project has a tool “f2py” that automates the process and hides the complications
- If you have a fortran file with some functions or subroutine do: `f2py -c code.f90 -m module`
 - Creates python loadable module “module”
 - Flag '-c' calls compiler; flag '-m' sets module name

Interfacing Fortran with f2py (2)

- Then in python do: `from module import *` and call the Fortran functions in python
- The f2py tool will parse the fortran code and generate the necessary C-code for a module
- The f2py generated code will automatically insert code to convert data as needed; e.g. lists are converted to arrays
- The f2py tool works best with well formed Fortran code; otherwise data maps can help

Fortran examples for f2py

- Example code that converts cleanly with f2py:

```
subroutine hello
  print*, "Hello, World!"
end subroutine hello
function sum_of_int(a,b) result(c)
  integer, intent(in) :: a, b
  integer :: c
  c = a + b
  print*, "sum of ", a, " and ", b, " is ", c
end function sum_of_int
function sum_of_double(a,b) result(c)
  double precision, intent(in) :: a, b
  double precision :: c
  c = a + b
  print*, "sum of ", a, " and ", b, " is ", c
end function sum_of_double
```

Passing arrays with f2py

- Arrays are traditional style arrays with f2py:

```
function sum_of_doubles(a,n) result(s)
  double precision, intent(in) :: a(*)
  integer, intent(in) :: n
  double precision :: s
  integer :: i
  s = 0
  do i=1,n
    s = s + a(i)
  end do
end function sum_of_doubles
```

```
-----
num = 10
dlist = [sqrt(float(i)) for i in range(1,num)]
dsum = sum_of_doubles(dlist,num)
```

Passing strings with f2py

- Strings are handled in a very similar fashion to traditional style arrays with f2py:

```
subroutine hello(name)
    character(len=*), intent(in) :: name
    print* , "Hello, ", name, "!"
end subroutine hello
```

```
-----
#!/usr/bin/env python
from hello import *
```

```
print ("Calling DS0")
hello('World')
print ("Done")
```

Compiling and Linking and Interfacing Multiple Programming Languages

Dr. Axel Kohlmeyer

Research Professor, Dept. of Mathematics
Associate Director, ICMS
College of Science and Technology
Temple University, Philadelphia

<http://sites.google.com/site/akohlmey/>

a.kohlmeyer@temple.edu